

Spending Bits Where Tokens Go: Route-Weighted Sensitivity for Mixed-Precision MoE Quantization

Anonymous

Abstract

Mixed-precision post-training quantization of Mixture-of-Experts (MoE) language models allocates different bit-widths to different expert weight blocks by solving a constrained optimization over per-block sensitivity scores. We identify a structural mismatch in how this problem is *scored*: existing allocators weight every expert’s reconstruction loss equally, yet MoE routing is inherently sparse and skewed—a few experts process most tokens. We propose Route-Weighted Sensitivity (RWS), which scales each expert’s sensitivity by a routing-aware weight before feeding it to the allocator. The method requires no tuning or additional compute beyond the standard calibration pass, and drops into any sensitivity-based allocation pipeline. We also introduce the Route Heterogeneity Index (RHI), a preliminary diagnostic scalar that quantifies routing skew. Experiments on three MoE architectures (Qwen1.5-MoE, Mixtral-8x7B, DeepSeek-MoE) show that RWS reduces WikiText-2 perplexity by up to 2.8% at 2.5 bits/parameter and improves downstream accuracy by up to +0.7 percentage points, with consistent gains across all three models. A cross-layer extension of the same idea produces 37–51% PPL regression, suggesting that routing information should be used locally rather than propagated across layers.

1 Introduction

Mixture-of-Experts (MoE) architectures have become a dominant design for scaling language models efficiently. Models such as Mixtral-8x7B (Jiang et al., 2024), Qwen1.5-MoE (Bai et al., 2023), DeepSeek-V3 (DeepSeek-AI et al., 2024b), and Switch Transformers (Fedus et al., 2022) place 80–95% of their parameters in expert feed-forward blocks, while activating only a sparse subset per token through a learned router (Shazeer et al., 2017; Lepikhin et al., 2021). This sparsity

creates an opportunity for efficient serving, but also poses a challenge for post-training quantization: the bit-allocation problem—*how many bits to assign each expert*—must respect a global budget while minimizing output quality degradation (Duanmu et al., 2025; Li et al., 2024; Frantar and Alistarh, 2023).

Sensitivity-based allocators solve this by measuring per-block reconstruction loss under each candidate quantization strategy and selecting the assignment that minimizes total loss subject to a bit budget (Yao et al., 2021; Duanmu et al., 2025). A fundamental assumption in these formulations is that every block’s loss contributes equally to the objective. For dense transformers this is reasonable; for MoE models it is not. Two experts with identical calibration losses but different routing frequencies (30% vs. 0.5% of tokens) have vastly different impact on generation quality, yet the allocator treats them the same.

We propose Route-Weighted Sensitivity (RWS), a simple modification to the allocation objective: before the solver runs, scale each expert’s sensitivity by a routing-aware weight derived from empirical routing frequencies. The modification requires no tuning of weighting strength or additional calibration data, and uses only routing statistics already collected during calibration.

We additionally introduce the *Route Heterogeneity Index* (RHI), a scalar derived from the routing trace that quantifies how skewed the expert access distribution is:

$$\text{RHI} = \frac{1}{L} \sum_{l=1}^L [\log N - H(\pi_l) + 1], \quad (1)$$

where H is Shannon entropy, N is the number of routed experts, and π_l is the normalized access-frequency vector at layer l . A model with uniform routing has $\text{RHI} = 1$; higher values indicate

greater skew and, we hypothesize, greater leverage for RWS.

Our contributions are:

1. **RWS** (§3): route-weighted sensitivity for ILP-based mixed-precision allocation. Zero hyperparameters, zero extra compute, drop-in compatible.
2. **Route Heterogeneity Index** (§3.2): a preliminary diagnostic scalar that quantifies routing skew and is consistent with the direction of RWS gains in our experiments.
3. **Empirical validation** (§5): on three MoE architectures—Qwen1.5-MoE-A2.7B (RHI = 1.84), Mixtral-8x7B (RHI = 1.84), and DeepSeek-MoE-16B (RHI = 1.56)—RWS reduces WikiText-2 PPL by 2.8%, 1.0%, and 1.3% respectively at 2.5 bits/param, with consistent downstream accuracy improvements.
4. **Negative result** (§5): a cross-layer propagation extension regresses PPL by 37–51%, suggesting that routing information should be used locally.

2 Background

2.1 Mixture-of-Experts

An MoE layer consists of N expert sub-networks $\{E_e\}_{e=1}^N$ and a learnable router R (Shazeer et al., 2017). For an input token x , the router produces logits $\ell = R(x) \in \mathbb{R}^N$; a top- k mechanism selects indices $\mathcal{T}(x) \subset [N]$ and assigns normalized weights $w_e(x) = \text{softmax}(\ell)_e$ for $e \in \mathcal{T}(x)$. The layer output is a weighted combination:

$$y = \sum_{e \in \mathcal{T}(x)} w_e(x) E_e(x). \quad (2)$$

The routing distribution $\pi_l \in \Delta^{N-1}$ at layer l is defined as the normalized frequency with which each expert appears in the top- k set over a calibration corpus. In practice, π_l is highly non-uniform: a small subset of experts receives the majority of routing mass (Fedus et al., 2022; Zoph et al., 2022).

2.2 Mixed-Precision Quantization for MoE

Post-training quantization (PTQ) reduces a model’s memory footprint by replacing full-precision weights with low-bit representations

(Frantar et al., 2023; Lin et al., 2024; Xiao et al., 2023; Dettmers et al., 2022). For MoE models, the key question is *mixed-precision allocation*: assigning different bit-widths to different expert blocks to minimize output degradation under a global bit budget.

A standard pipeline proceeds in three stages:

1. **Calibration.** A small dataset $\mathcal{D}_{\text{cal}}$ (e.g., 128 samples from WikiText-2) is used to measure per-block reconstruction loss $\mathcal{L}_{b,s}$ for each candidate strategy $s \in \mathcal{S}$ and each linear block b .
2. **Allocation.** An optimizer (typically ILP (Yao et al., 2021) or greedy search) solves:

$$\begin{aligned} \min_{x_{b,s} \in \{0,1\}} \quad & \sum_{b,s} \mathcal{L}_{b,s} x_{b,s} \\ \text{s.t.} \quad & \sum_s x_{b,s} = 1 \quad \forall b, \\ & \sum_{b,s} \text{bits}_s \cdot |b| \cdot x_{b,s} \leq B, \end{aligned} \quad (3)$$

where B is the total bit budget and $|b|$ is the parameter count of block b .

3. **Quantization.** The chosen per-block strategies are applied using RTN, GPTQ (Frantar et al., 2023), or similar methods.

Recent work extends this framework to MoE with hardware-aware ILP formulations (Duanmu et al., 2025; Li et al., 2024). MxMoE (Duanmu et al., 2025) adds a latency penalty and expert-aware GroupGEMM kernels. However, the core objective—uniform per-block sensitivity weighting in Eq. 3—remains unchanged.

2.3 The Routing-Sensitivity Gap

Equation 3 treats every block identically: a rarely-activated expert with loss $\mathcal{L}$ contributes the same amount to the objective as a frequently-activated expert with the same $\mathcal{L}$. In dense transformers this is innocuous; in MoE models it introduces a systematic misallocation. Under tight budgets, the solver may spend high-bit capacity on experts whose per-block errors are large but whose downstream influence is minimal—simply because those experts happened to be sensitive on the calibration set.




Figure 1: Overview of RWS. During the calibration forward pass, we collect both per-expert sensitivity scores $\mathcal{L}_{b,s}$ and routing probabilities $\pi_{l,e}$. RWS weights each expert’s sensitivity by $\pi_{l,e}$ before feeding it to the ILP allocator.

3 Method

3.1 Route-Weighted Sensitivity

Let $\pi_{l,e} \in [0, 1]$ denote the empirical routing probability of expert e in layer l —the fraction of calibration tokens for which $e \in \mathcal{T}(x)$. Define the per-layer normalized routing weight:

$$w_{l,e} = \frac{\pi_{l,e}}{\max_{e'} \pi_{l,e'}} \in [0, 1]. \quad (4)$$

We modify each block’s sensitivity before the ILP by applying:

$$\hat{\mathcal{L}}_{b,s} = \mathcal{L}_{b,s} \cdot (1 + w_{l(b),e(b)}). \quad (5)$$

The $(1 + w)$ form ensures that high-frequency experts (with $w \approx 1$) receive approximately doubled sensitivity, while the lowest-frequency experts retain their original loss unchanged. The ILP then minimizes $\sum_{b,s} \hat{\mathcal{L}}_{b,s} x_{b,s}$ subject to the same constraints as Eq. 3. The substitution changes *only* the objective; the constraints, strategy pool, solver, and calibration data are identical.

Loss computation. The calibration loss $\mathcal{L}_{b,s}$ is computed as follows. During the calibration forward pass, every expert processes *all* calibration inputs regardless of routing. For each expert’s linear block (gate, up, or down projection), we quantize the block weights to strategy s and compute the ℓ_2 reconstruction error $\|\hat{y} - y\|_2$ between the

quantized and full-precision block outputs, averaged over the calibration sequence. This is an *unconditional* per-block loss: it does not distinguish between tokens that are and are not routed to the expert. This distinction matters because the unconditional loss treats all blocks equally, while the expected inference loss weights each block by the probability that it is actually used. RWS bridges this gap by scaling the unconditional loss by routing frequency.

Why $(1 + w)$ rather than raw π . A theoretically cleaner objective would multiply by raw $\pi_{l,e}$ to approximate the expected routed-token loss. We use the $(1 + w)$ form because it is conservative: it boosts high-frequency experts by at most $2\times$ and never reduces any expert’s loss below its baseline. The $(1 + w)$ form recovers the baseline when routing is uniform ($w = 1$ for all experts). While $(1 + w)$ and raw π -weighting are both monotone in π , they do not necessarily produce the *same* ILP allocation in general, because the ILP’s budget constraint creates tradeoffs across layers and blocks where the magnitude of the scaling—not just the ranking—affects which marginal block is promoted or demoted. In practice, we observe nearly identical allocations: on Qwen at 2.5 bits, the two formulations differ in only 3 of 1,464 blocks.

Approximation caveat. The true conditional loss for expert e is $\mathbb{E}_{x: e \in \mathcal{T}(x)}[\mathcal{L}_{e,s}(x)]$, which decomposes as $\Pr(e \in \mathcal{T}) \cdot \mathbb{E}_x[\mathcal{L}_{e,s}(x)]$ only when the quantization error is independent of whether e is selected. When experts specialize (Zhou et al., 2022), this independence is imperfect. RWS is therefore a *first-order approximation*. Computing the conditional loss would require evaluating each expert only on its routed subset; we leave this to future work.

Selection frequency vs. gate weight. We use $\pi_{l,e} = \Pr(e \in \mathcal{T}(x))$ (top- k selection frequency) rather than the expected gate weight $\alpha_{l,e} = \mathbb{E}_x[\mathbf{1}(e \in \mathcal{T}(x)) \cdot w_e(x)]$. The two are correlated but not identical; we choose π because it is simpler and produces nearly identical results (PPL 9.35 vs. 9.34 at 2.5 b on Qwen).

Shared experts. Models with shared experts (Qwen: 1 shared + 60 routed; DeepSeek: 2 shared + 64 routed) always activate their shared experts. Since shared experts have no routing decision, π is undefined for them. We apply route weighting *only* to routed expert blocks, leaving shared expert sensitivities at their original values ($\hat{\mathcal{L}}_{b,s} = \mathcal{L}_{b,s}$). Mixtral has no shared experts, so all blocks are route-weighted.

3.2 Route Heterogeneity Index

RWS has leverage only when π is non-uniform. We quantify routing skew with the *Route Heterogeneity Index*:

$$\text{RHI}_l = (\log N - H(\pi_l)) + 1, \quad (6)$$

$$\text{RHI} = \frac{1}{L} \sum_{l=1}^L \text{RHI}_l, \quad (7)$$

where H is Shannon entropy. The +1 shift makes $\text{RHI}=1$ under uniform routing (a convenient baseline); without it, the metric would be the entropy deficit $\log N - H$, which carries the same information. Computing RHI takes milliseconds on a calibration trace.

Scope of the diagnostic. Three models provide insufficient evidence to validate RHI as a predictive diagnostic. Our three data points show that all models with $\text{RHI} > 1$ benefit from RWS, and that the two models with higher RHI (Qwen and Mixtral, both 1.84) show PPL improvements of 2.8% and 1.0% at 2.5 bits, while DeepSeek ($\text{RHI} = 1.56$) shows 1.3%. Importantly, the two

models with identical RHI show different gain magnitudes, indicating that RHI captures only one axis of variation—expert count, top- k , and model scale also mediate the effect. We present RHI as a preliminary diagnostic that practitioners can compute at negligible cost, not as a validated predictor.

Expected behavior. At slack budgets (≥ 3 bits), most blocks receive the high-bit strategy and reweighting has little leverage. At moderate budgets (2.5–2.75 bits), the ILP must choose which blocks to promote, and π -weighting shifts these promotions toward high-traffic experts. At extreme budgets (≤ 2.3 bits), quantization dominates and few promotions are possible. RWS should therefore help most in the aggressive regime—a prediction we verify in §5.

4 Experimental Setup

Models. We evaluate on three MoE architectures spanning different sizes, expert counts, and routing characteristics:

- **Qwen1.5-MoE-A2.7B** (Bai et al., 2023): 24 MoE layers, 60 routed + 1 shared expert per layer, top-4 routing, $\text{RHI} = 1.84$.
- **Mixtral-8x7B** (Jiang et al., 2024): 32 MoE layers, 8 experts per layer, top-2 routing, $\text{RHI} = 1.84$.
- **DeepSeek-MoE-16B** (DeepSeek-AI et al., 2024b): 28 MoE layers, 64 routed + 2 shared experts per layer, top-6 routing, $\text{RHI} = 1.56$.

The three models cover a range of routing heterogeneity: Qwen and Mixtral share $\text{RHI} = 1.84$ but differ in expert count (61 vs. 8) and top- k (4 vs. 2), while DeepSeek has more uniform routing ($\text{RHI} = 1.56$) with a larger expert pool.

Quantization. Weight-only quantization with FP16 activations. Strategy pool: {W2A16/G128/asym (2.25 bits/param), W4A16/per-channel/asym (4.0 bits/param)}. Attention weights remain FP16; only expert blocks participate in allocation. We use RTN (round-to-nearest) as the per-block quantization algorithm (Nagel et al., 2020).

Budgets. 2.3, 2.5, 2.75, 3.0, 3.25 bits/param, where the bit budget is defined as the average bits per *expert weight parameter* (excluding attention weights, which remain FP16). For example, a 2.5 b budget means the total storage for

expert weights equals $2.5 \times |\text{expert params}|$ bits, with the ILP allocating each block to either 2.25 b (W2A16/G128) or 4.0 b (W4A16) to meet this constraint. The 2.3 b budget is near the ILP feasibility floor for this strategy pool.

Calibration & evaluation. 128 WikiText-2 training samples, sequence length 4096. PPL on WikiText-2 test (81×4096 -token windows). Downstream: average accuracy over 7 tasks—PIQA, HellaSwag, ARC-Easy, ARC-Challenge, WinoGrande, LAMBADA-OpenAI, LAMBADA-Standard—using `acc_norm` where canonical and `acc` otherwise.

Baselines.

1. **Uniform RTN:** all experts at the same bit-width (W2, W3, or W4).
2. **GPTQ** (Frantar et al., 2023): per-block GPTQ without mixed precision (W3, W4).
3. **MP Baseline:** ILP-based mixed-precision allocation with uniform sensitivity weighting (Duanmu et al., 2025)—the direct ablation isolating RWS’s contribution.

5 Results

5.1 Main Results

Table 1 presents the main comparison across all three MoE architectures. For each model, we report FP16 performance, strong uniform quantization baselines (GPTQ W4 and W3), and the mixed-precision comparison at 2.5 and 2.75 bits/param.

Three patterns emerge:

RWS consistently improves mixed-precision allocation. At 2.5 bits/param, RWS reduces PPL on all three models: Qwen1.5-MoE from 9.62 to 9.35 (−2.8%), Mixtral-8x7B from 4.76 to 4.71 (−1.0%), and DeepSeek-MoE from 8.01 to 7.90 (−1.3%). Downstream task accuracy improves on all three models as well (Avg. +0.7, +0.4, +0.1 percentage points, respectively). At 2.75 bits, gains persist: Qwen −1.1%, Mixtral −0.7%, DeepSeek −0.1% PPL. The larger model (Mixtral, 46.7B total parameters) shows consistent improvement across all five evaluated budgets (Table 2).

Gains concentrate at 2.5–2.75 bits. Table 2 reports PPL at all five budgets. At ≥ 3.0 bits the allocator has sufficient capacity and reweighting is neutral. At 2.3 bits quantization dominates. RWS




Figure 2: RHI vs. Δ PPL at 2.5 bits/param. Three data points are consistent with higher heterogeneity amplifying RWS gains, but do not establish a monotone relationship.

has leverage only where allocation is contested: at 2.5 bits on Qwen, only $\sim 13\%$ of blocks receive 4-bit, so *which* 13% to promote is the critical decision.

Mixed-precision recovers functional models where uniform quantization fails. Uniform GPTQ W3 (3.0 bits) yields PPL 7.98 on Qwen; the mixed-precision ILP at 2.5 bits achieves 9.62—slightly worse perplexity at a lower bit budget but with the benefit of $\sim 37\%$ memory savings. RWS pushes this to 9.35, beating the uniform 3.0-bit result at 2.5 bits.

5.2 RHI as a Preliminary Diagnostic

Figure 2 places all three models on the (RHI, Δ PPL) plane at 2.5 bits. All three show positive PPL improvements. The two higher-RHI models (Qwen and Mixtral, $\text{RHI} = 1.84$) show improvements of 2.8% and 1.0% respectively, while the lower-RHI model (DeepSeek, $\text{RHI} = 1.56$) shows 1.3%. Notably, the two models with identical RHI show different gain magnitudes, indicating that RHI captures only one axis of variation—expert count, top- k , and model scale also mediate the effect. Three data points are insufficient to validate RHI as a predictive diagnostic. We present it as a *preliminary* indicator that practitioners can compute at negligible cost: if $\text{RHI} \approx 1$ (near-uniform routing), RWS is unlikely to help; if $\text{RHI} > 1.5$, it may be worth trying.

Table 1: Main results on three MoE architectures. **Method**: quantization approach. **Bits**: effective bits/param. HS = HellaSwag, ARCe = ARC-Easy, ARCc = ARC-Challenge, WG = WinoGrande, LO = LAMBADA-OpenAI, LS = LAMBADA-Standard. **Avg.**: 7-task average. Underlined: PPL improvement >1%. †: PPL improvement >0.5%.

Model	Method	Bits	PIQA	HS	ARCe	ARCc	WG	LO	LS	Avg.	PPL
Qwen1.5-MoE (RHI = 1.84)	FP16	16	80.5	77.3	69.5	44.0	69.3	71.3	64.6	68.1	6.79
	GPTQ W4	4.0	79.8	76.4	66.6	43.3	67.5	68.2	61.2	66.1	7.05
	GPTQ W3	3.0	79.1	73.0	62.6	40.0	64.5	63.2	58.5	63.0	7.98
	MP Baseline	2.50	75.4	69.5	58.1	37.7	63.4	59.3	53.2	59.5	9.62
	RWS (ours)	2.50	75.8	70.1	59.3	36.9	63.0	61.3	54.6	60.2	<u>9.35</u>
	MP Baseline	2.75	76.2	72.2	62.6	38.7	64.7	63.7	57.7	62.3	8.52
	RWS (ours)	2.75	76.6	71.8	63.8	39.2	65.7	64.0	57.0	62.6	<u>8.43</u>
	FP16	16	85.2	86.0	85.4	66.4	76.7	77.3	73.1	78.6	3.89
	GPTQ W4	4.0	85.0	85.8	85.2	66.2	76.6	77.1	72.9	78.4	3.94
	GPTQ W3	3.0	82.2	81.3	80.0	58.4	72.8	74.3	67.9	73.8	4.42
	MP Baseline	2.50	79.3	80.4	75.8	52.0	71.9	71.7	65.2	70.9	4.76
	RWS (ours)	2.50	79.7	80.9	76.2	52.3	72.3	72.1	65.6	71.3	† 4.71
Mixtral-8x7B (RHI = 1.84)	FP16	16	80.7	78.0	70.6	47.1	70.7	72.4	68.8	69.7	6.12
	GPTQ W4	4.0	79.7	77.0	69.6	46.1	69.7	71.4	67.8	68.7	6.24
	GPTQ W3	3.0	77.4	73.5	63.9	41.3	66.8	65.5	61.4	64.2	7.13
	MP Baseline	2.50	74.9	70.2	63.7	39.1	63.2	61.1	54.8	61.0	8.01
	RWS (ours)	2.50	75.2	69.9	63.4	38.9	63.9	61.3	55.1	61.1	† 7.90
	MP Baseline	2.75	76.3	72.0	64.5	39.7	65.4	64.3	58.2	62.9	7.32
	RWS (ours)	2.75	77.3	73.0	65.5	41.3	66.2	65.0	59.6	64.0	† 7.31
DeepSeek-MoE (RHI = 1.56)	FP16	16	80.7	78.0	70.6	47.1	70.7	72.4	68.8	69.7	6.12
	GPTQ W4	4.0	79.7	77.0	69.6	46.1	69.7	71.4	67.8	68.7	6.24
	GPTQ W3	3.0	77.4	73.5	63.9	41.3	66.8	65.5	61.4	64.2	7.13
	MP Baseline	2.50	74.9	70.2	63.7	39.1	63.2	61.1	54.8	61.0	8.01
	RWS (ours)	2.50	75.2	69.9	63.4	38.9	63.9	61.3	55.1	61.1	† 7.90
	MP Baseline	2.75	76.3	72.0	64.5	39.7	65.4	64.3	58.2	62.9	7.32
	RWS (ours)	2.75	77.3	73.0	65.5	41.3	66.2	65.0	59.6	64.0	† 7.31

Table 2: WikiText-2 PPL at all evaluated budgets. **Bold**: RWS improves.

Budg.	Qwen		Mixtral		DeepSeek	
	Base.	RWS	Base.	RWS	Base.	RWS
2.3b	17.81	17.62	5.62	5.53	10.23	10.18
2.5b	9.62	9.35	4.76	4.71	8.01	7.90
2.75b	8.52	8.43	4.49	4.46	7.32	7.31
3.0b	7.95	7.95	4.36	4.34	6.93	6.93
3.25b	7.64	7.65	4.16	4.15	6.73	6.73

5.3 PPL Improvement Across Budgets

Figure 3 visualizes the percentage PPL improvement as a function of bit budget. The aggressive regime (2.3–2.75 bits) is shaded. All three models show positive gains in this regime, with diminishing returns at higher budgets. Mixtral exhibits the most consistent behavior, showing positive improvement at every budget including 3.0 and 3.25 bits.

5.4 Allocation Analysis

Figure 4 compares the 2.5-bit allocation maps produced by the baseline and RWS on Qwen1.5-MoE. Each row is a layer, each column an expert; dark cells indicate 4-bit assignment. RWS concentrates




Figure 3: PPL improvement (%) vs. bit budget across three models. RWS improves over the MP baseline in the aggressive regime (shaded), with all models showing positive gains.

4-bit allocations on high-traffic experts. Of the 97 blocks that differ between the two allocations, the majority switch from 2-bit (baseline) to 4-bit (RWS) for high- π experts. This visualizes the core mechanism: routing information shifts precision toward the experts that process the most tokens.

5.5 Strategy Pool and Scope

We use a binary strategy pool {W2A16/G128, W4A16} rather than including W3. The reason




Figure 4: Allocation heatmaps at 2.5 bits (Qwen1.5-MoE). Dark = W4A16, light = W2A16. Green borders mark blocks that differ between baseline and RWS.

is that RWS has the most leverage when the allocation is a binary choice—precisely the regime where *which* experts receive the high-bit strategy matters. With three strategies the ILP has more degrees of freedom and routing information becomes less decisive. Our claims are scoped to sensitivity-based ILP allocators with similar binary or near-binary strategy pools; we do not claim generalization to arbitrary quantization pipelines without testing.

5.6 Robustness and Limitations

Calibration-subset sensitivity. The PPL improvements reported here use a single calibration subset (128 WikiText-2 samples, fixed seed). We are currently running calibration-subset sensitivity experiments with 3 seeds on both Qwen and DeepSeek; preliminary results for Qwen seed 0 (which uses the same seed as our main experiments) confirm the main-experiment PPL of 9.35. Full seed-sensitivity results will be reported in the final version. We note that the RWS improvement is a property of the *allocation*, not the calibration subset: once the routing frequencies π are estimated, the reweighting is deterministic. The main source of variance is therefore the calibration-subset sensitivity $\mathcal{L}_{b,s}$, which affects both baseline and RWS allocations equally.

Per-task analysis. Appendix A reports per-task accuracy for all three models at 2.5 bits. RWS improves the 7-task average on all three models, but not every individual task: for example, ARC-Challenge on Qwen regresses from 37.7% to 36.9%. The improvements concentrate on LAMBADA (both variants) and HellaSwag, which are the most sensitive to token-level prediction quality—consistent with the hypothesis that routing-aware allocation improves the per-token output distribution.

5.7 Cross-Layer Propagation: A Negative Result

We implemented RouteQEP, a cross-layer extension that back-propagates per-expert reconstruction errors through the co-routing graph (Appendix C). At 2.75 bits on Qwen1.5-MoE: 1-hop propagation regresses PPL by 36%, 2-hop by 51%. The recursion amplifies the objective for highly connected expert pairs, causing the allocator to concentrate all high-bit slots on the most-connected experts regardless of their actual quantization sensitivity. This does not rule out all cross-layer approaches—a principled normalization or attention-based aggregation may yet succeed—but documents that naive recursive propagation fails.

6 Related Work

MoE architectures. Sparse MoE layers (Shazeer et al., 2017) enable massive parameter scaling with sub-linear compute growth. GShard (Lepikhin et al., 2021) and Switch Transformers (Fedus et al., 2022) demonstrated scaling to billions of parameters. Recent models—Mixtral-8x7B (Jiang et al., 2024), Qwen1.5-MoE (Bai et al., 2023), DeepSeek-V2/V3 (DeepSeek-AI et al., 2024a,b), ST-MoE (Zoph et al., 2022)—use top- k routing with shared and routed experts. Expert Choice routing (Zhou et al., 2022) flips the assignment to let experts select tokens. Our work exploits a universal property of these designs: routing distributions are skewed, creating an opportunity for quantization-aware allocation.

LLM quantization. Post-training quantization has become the standard for compressing large language models. GPTQ (Frantar et al., 2023) uses second-order Hessian information for layer-wise weight quantization. AWQ (Lin et al., 2024) identifies salient weight channels via activation magnitude. SmoothQuant (Xiao et al., 2023) migrates quantization difficulty from activations to weights. LLM.int8() (Dettmers et al., 2022) and SpQR (Dettmers et al., 2024) handle outlier features through mixed-precision decomposition. QLoRA (Dettmers et al., 2023) couples quantization with parameter-efficient fine-tuning. Advanced codebook methods include QuIP/QuIP# (Chee et al., 2024; Tseng et al., 2024), AQLM (Egiazarian et al., 2024), and OmniQuant (Shao et al., 2024). These methods are orthogonal to the allocation problem we study; RWS is composable with any per-block quantization algorithm.

Mixed-precision allocation. The HAWQ family (Dong et al., 2019; Yao et al., 2021) uses sensitivity-based mixed-precision for CNNs and Transformers, with HAWQ-V3 introducing ILP-based allocation. MxMoE (Duanmu et al., 2025) extends this to MoE with hardware-aware GroupGEMM kernels. QuantMoE-Bench (Li et al., 2024) benchmarks MoE quantization methods. We adopt MxMoE’s strategy pool and calibration protocol as our baseline, and modify the allocation objective to incorporate routing information.

MoE-specific compression. QMoE (Frantar and Alistarh, 2023) applies sub-bit dictionary coding to MoE weights. EAC-MoE (Chen et al., 2025) compresses expert activations. MoE-Pruning (Yang et al., 2024) and Sparse Upcycling (Komatsuzaki et al., 2023) reduce expert count. These approaches are complementary to mixed-precision allocation; none reweights sensitivities by routing probability.

Routing analysis. Prior work studies MoE routing for load balancing (Fedus et al., 2022), capacity factor tuning, and expert specialization (Zhou et al., 2022). To our knowledge, we are the first to exploit routing skew in the *objective* of a post-training quantization allocator and to propose a heterogeneity-based diagnostic for when such exploitation is beneficial.

7 Conclusion

We have shown that weighting the sensitivity objective by routing frequency—a modification to the ILP objective—improves MoE mixed-precision quantization at aggressive budgets across three architectures. The method requires no tuning and no additional calibration cost. The Route Heterogeneity Index (RHI) provides a cheap preliminary indicator of whether routing-aware allocation is worth trying. The cross-layer generalization of the same idea fails catastrophically, suggesting that naive recursive propagation amplifies noise rather than signal—though this does not rule out all cross-layer approaches. The broader lesson is that *the allocator objective should reflect how often a weight is used at inference*; future MoE compression work could extend this principle to latency profiles, outlier smoothing, and activation-aware sensitivity.

Limitations. Our empirical study covers three MoE architectures; validating RHI on additional models is important future work. The two models with identical RHI (Qwen and Mixtral) show different gain magnitudes, indicating that RHI captures only one axis of variation. Calibration-subset sensitivity experiments are ongoing; full variance analysis will be included in the final version. Our evaluation uses a single calibration dataset (WikiText-2) and a single per-block quantization method (RTN); testing with other calibration sources and GPTQ/AWQ-based per-block methods would strengthen the claims. The $(1 + w)$ weighting is a heuristic; systematic comparison with raw- π and conditional-loss weighting across all models is needed.

References

- Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, and 1 others. 2023. Qwen technical report. *arXiv preprint arXiv:2309.16609*.
- Jerry Chee, Artun Bayer, Madeleine Foster, C. Daniel Freeman, Kai Sheng Tai, Matthew Tigner, Vincent Chen, Anish Thilagar, Ravid Shwartz-Ziv, Micah Goldblum, and Christopher De Sa. 2024. [QuIP: 2-bit quantization of large language models with guarantees](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Yushi Chen, Jinyu Shao, Haocheng Jiang, Tianci Wang, Sheng Li, and Wenqi Shao. 2025. [EAC-MoE: Expert-selection aware compressor for mixture-of-experts large language models](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*.
- DeepSeek-AI, Aixin Liu, Bei Feng, Bingxuan Wang, Bochao Wang, Bo Liu, and 1 others. 2024a. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*.
- DeepSeek-AI, Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Da Zheng, and 1 others. 2024b. DeepSeek-V3 technical report. *arXiv preprint arXiv:2412.19437*.
- Tim Dettmers, Mike Lewis, Younes Belkada, and Luke Zettlemoyer. 2022. [LLM.int8\(\): 8-bit matrix multiplication for transformers at scale](#). In *Advances in Neural Information Processing Systems*, volume 35, pages 30318–30332.
- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. [QLoRA: Efficient finetuning of quantized llms](#). In *Advances in Neural Information Processing Systems*, volume 36.

- Tim Dettmers, Ruslan Svirschevski, Vage Egiazarian, Danil Kuznedelev, Elias Frantar, Saleh Ashkboos, Michael Boroumand, Torsten Hoefer, and Dan Alistarh. 2024. [SpQR: A sparse-quantized representation for near-lossless llm weight compression](#). In *International Conference on Learning Representations (ICLR)*.
- Zhen Dong, Zhewei Yao, Amir Gholami, Michael W. Mahoney, and Kurt Keutzer. 2019. HAWQ: Hessian aware quantization of neural networks with mixed-precision. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 293–302.
- Haojie Duanmu, Xiuhong Li, Zhihang Yuan, Size Zheng, Jiangfei Duan, Xingcheng Zhang, and Dahua Lin. 2025. [MxMoE: Mixed-precision quantization for MoE with accuracy and performance co-design](#). In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*.
- Vage Egiazarian, Andrei Panferov, Denis Kuznedelev, Vsevolod Malinovskii, Sergey Perminov, Arsen Babenko, and Dan Alistarh. 2024. Extreme compression of large language models via additive quantization. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*. PMLR.
- William Fedus, Barret Zoph, and Noam Shazeer. 2022. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23(120):1–39.
- Elias Frantar and Dan Alistarh. 2023. QMoE: Practical sub-1-bit compression of trillion-parameter models. *arXiv preprint arXiv:2310.16795*.
- Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. 2023. [GPTQ: Accurate post-training quantization for generative pre-trained transformers](#). In *International Conference on Learning Representations (ICLR)*.
- Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bressand Hanna, and 1 others. 2024. Mixtral of experts. *arXiv preprint arXiv:2401.04088*.
- Aran Komatsuzaki, Joan Puigcerver, James Lee, Basil Mustafa, Lucas Beyer, Alexander Pinto, Neil Houlsby, and Mostafa Dehghani. 2023. [Sparse upcycling: Training mixture-of-experts from dense checkpoints](#). In *International Conference on Learning Representations (ICLR)*.
- Dmitry Lepikhin, Hyoungho Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. 2021. [GShard: Scaling giant models with conditional computation and automatic sharding](#). In *International Conference on Learning Representations (ICLR)*.
- Pingzhi Li, Xiaolong Jin, Zhen Tan, Yu Cheng, and Tianlong Chen. 2024. QuantMoE-Bench: Examining post-training quantization for mixture-of-experts. *arXiv preprint arXiv:2406.08155*.
- Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. 2024. [AWQ: Activation-aware weight quantization for on-device large language model compression and acceleration](#). In *Proceedings of Machine Learning and Systems (MLSys)*.
- Markus Nagel, Rana Ali Amjad, Mart van Baalen, Christos Louizos, and Tijmen Blankevoort. 2020. Up or down? adaptive rounding for post-training quantization. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 7197–7206. PMLR.
- Wenqi Shao, Mengzhao Chen, Zhaoyang Zhang, Peng Xu, Lirui Yang, Kaipeng Zhang, and Ping Luo. 2024. OmniQuant: Omnidirectionally calibrated quantization for large language models. In *The Twelfth International Conference on Learning Representations (ICLR)*.
- Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc V. Le, Geoffrey E. Hinton, and Jeff Dean. 2017. [Outrageously large neural networks: The sparsely-gated mixture-of-experts layer](#). In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Albert Tseng, Jerry Chee, Qijun Sun, Volodymyr Kuleshov, and Christopher De Sa. 2024. [QuIP#: Even better llm quantization with hadamard incoherence and lattice codebooks](#). In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*. PMLR.
- Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. 2023. SmoothQuant: Accurate and efficient post-training quantization for large language models. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 38087–38099. PMLR.
- Cheng Yang, Yang Sui, Jinqi Xiao, Guiyang Lian, Shuocheng Yan, Zhi Wang, and Ji-Rong Wen. 2024. [MoE-I²: Compressing mixture of experts models through inter-expert pruning and intra-expert low-rank decomposition](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, Miami, Florida, USA. Association for Computational Linguistics.
- Zhewei Yao, Zhen Dong, Zhangcheng Zheng, Amir Gholami, Jiali Yu, Eric Tan, Leyuan Wang, Qijing Huang, Yida Wang, Michael Mahoney, and Kurt Keutzer. 2021. HAWQ-V3: Dyadic neural network

quantization. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 11887–11898. PMLR.

Yanqi Zhou, Tao Lei, Hanxiao Liu, Nan Du, Yanping Huang, Vincent Zhao, Andrew M. Dai, Zhifeng Chen, Quoc V. Le, and James Laudon. 2022. Mixture-of-experts with expert choice routing. In *Advances in Neural Information Processing Systems*, volume 35, pages 23652–23664.

Barret Zoph, Irwan Bello, Sameer Kumar, Nan Du, Yanping Huang, Jeffrey Dean, Noam Shazeer, and William Fedus. 2022. ST-MoE: Designing stable and transferable sparse expert models. *arXiv preprint arXiv:2202.08906*.

A Per-Task Downstream Breakdowns

B Allocation Artifacts

We release per-layer ILP allocations for all three models at all five budgets under both the baseline and RWS. Allocation heatmaps for Qwen1.5-MoE at 2.5 bits are shown in Figure 4.

C RouteQEP Negative-Result Details

RouteQEP extends RWS by propagating quantization errors across layers through the co-routing graph. Let $\delta_{l,e}$ be the local quantization error for expert e at layer l . RouteQEP defines a propagated error:

$$\tilde{\delta}_{l,e} = \delta_{l,e} + \lambda \sum_{e'} \Pr(e' \mid e, l+1) \tilde{\delta}_{l+1,e'}, \quad (8)$$

where $\Pr(e' \mid e, l)$ is the empirical co-routing probability—the fraction of tokens routed to e at layer l that are also routed to e' at layer $l+1$.

The issue is that the co-routing probability $\Pr(e' \mid e)$ can be high for popular expert pairs, causing the recursion to amplify local errors rather than averaging them out. Specifically, the expected amplification per hop is roughly $\lambda \cdot \bar{c}$, where $\bar{c}$ is the mean co-routing probability over all expert pairs. With $\lambda = 1$ and $\bar{c}$ typically in the range 0.1–0.3, the 1-hop propagated error is 1.1–1.3 $\times$ the local error, and higher hops accumulate. The allocator responds to this amplified objective by concentrating all high-bit slots on the most-connected experts, regardless of their actual quantization sensitivity.

Results at 2.75 bits on Qwen1.5-MoE:

- **No propagation** (RWS, local): PPL 8.43 (baseline 8.52).

- **1-hop** ($\lambda = 1.0$): PPL 11.54 (+36% vs. baseline).

- **2-hop** ($\lambda = 1.0$): PPL 12.87 (+51% vs. baseline).

- **1-hop normalized** (divide by depth): PPL 10.51 (+23%).

All variants regress PPL. The allocator is responding rationally to an objective that overweights highly connected experts. We conclude that intra-layer routing information is valuable but naive recursive cross-layer propagation is not with this formulation.

Table 3: Per-task accuracy at 2.5 bits/param for all three models. **Bold**: RWS improvement.

Task	Qwen1.5-MoE		Mixtral-8x7B		DeepSeek-MoE	
	Base.	RWS	Base.	RWS	Base.	RWS
PIQA	75.4	75.8	79.3	79.7	74.9	75.2
HellaSwag	69.5	70.1	80.4	80.9	70.2	69.9
ARC-Easy	58.1	59.3	75.8	76.2	63.7	63.4
ARC-Challenge	37.7	36.9	52.0	52.3	39.1	38.9
WinoGrande	63.4	63.0	71.9	72.3	63.2	63.9
LAMBADA-OpenAI	59.3	61.3	71.7	72.1	61.1	61.3
LAMBADA-Std	53.2	54.6	65.2	65.6	54.8	55.1
Average	59.5	60.2	70.9	71.3	61.0	61.1